Multilabels Classification of Podcast Genres

An Experiment with Multimodal Neural Networks

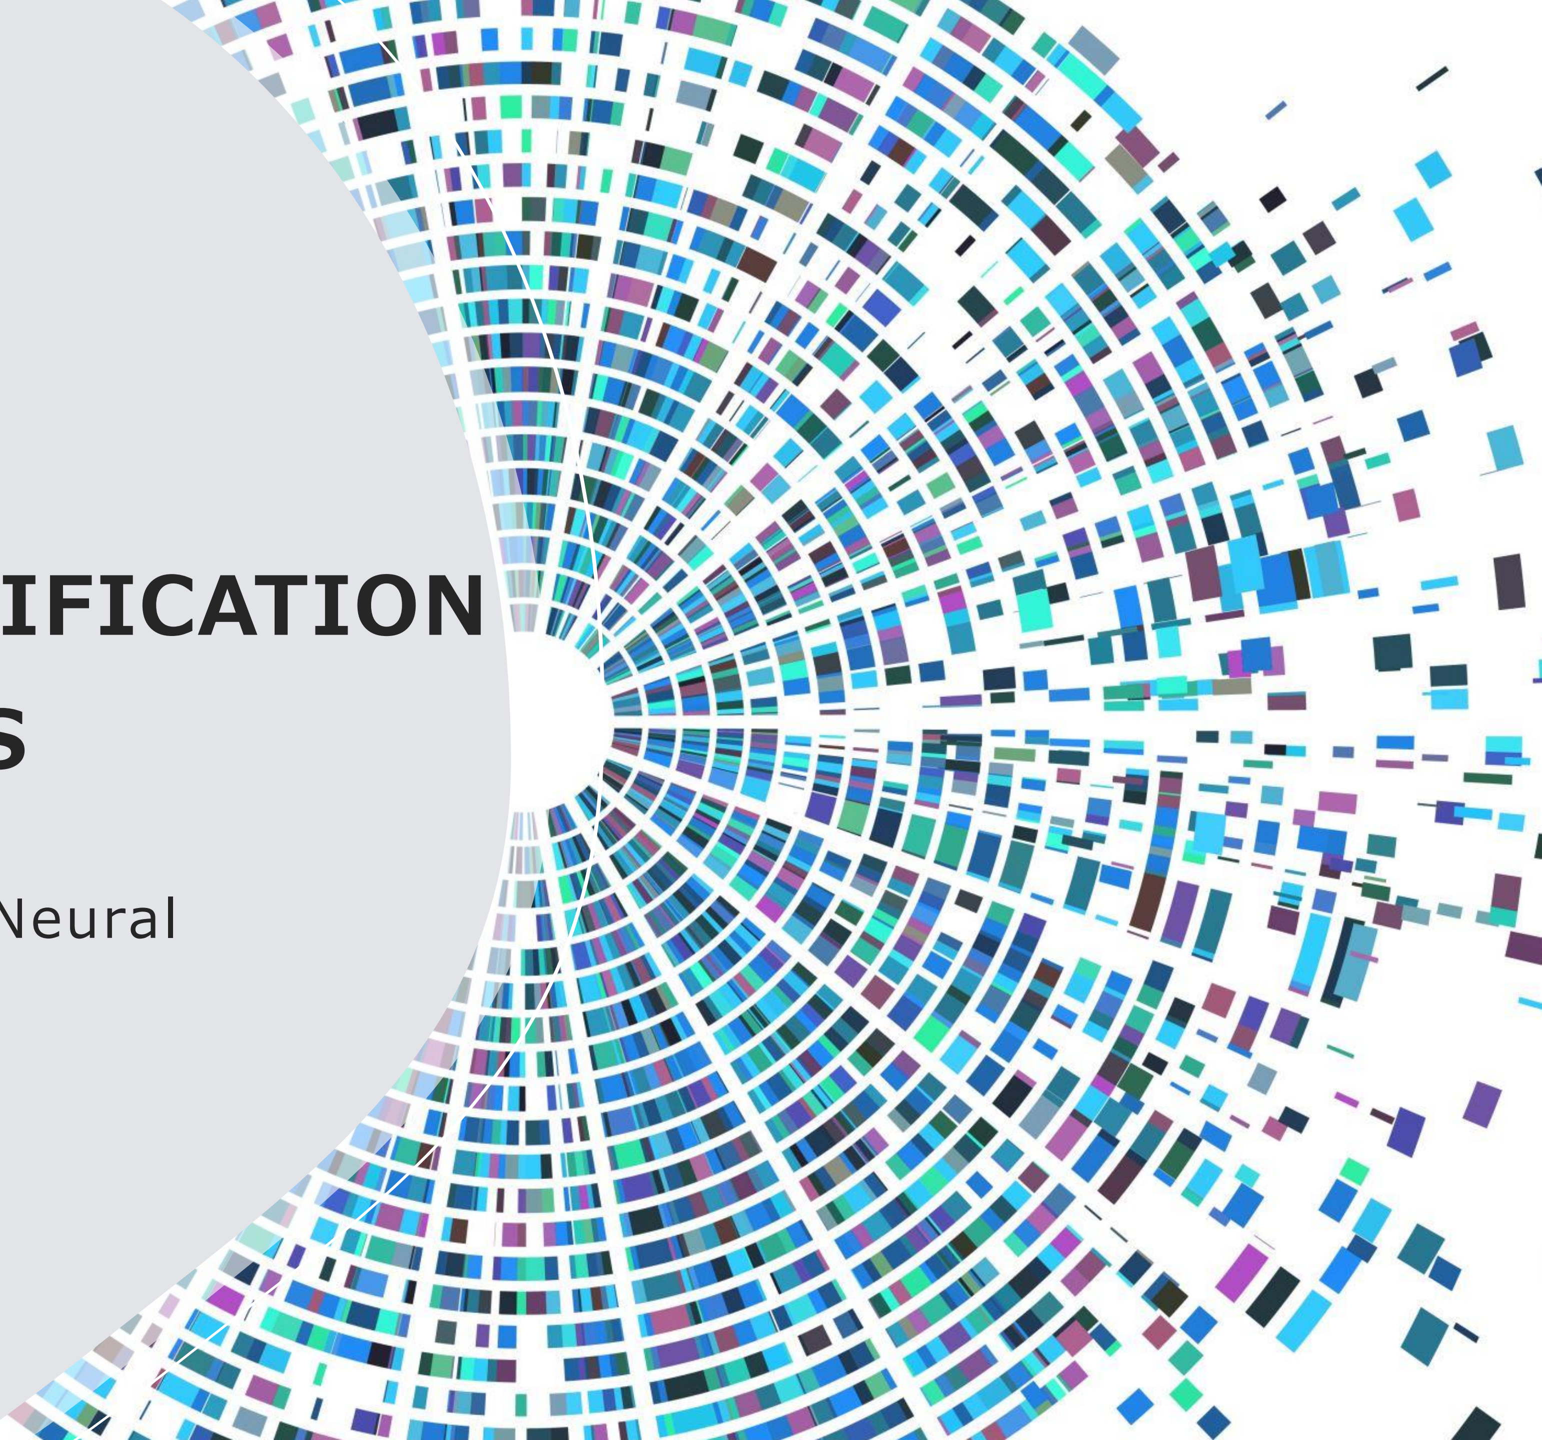

Podcasts

Podcasts are an emerging medium on the internet and the number of shows is growing exponentially over the years

They provide huge amounts of data for many different tasks (e.g. audio transcription)

Streaming services (e.g. Spotify, Apple Music) usually classify shows on the basis of manual annotations and use those labels to suggest new contents to their users

Is it possible to automatically classify podcasts according to the metadata available?

Project workflow

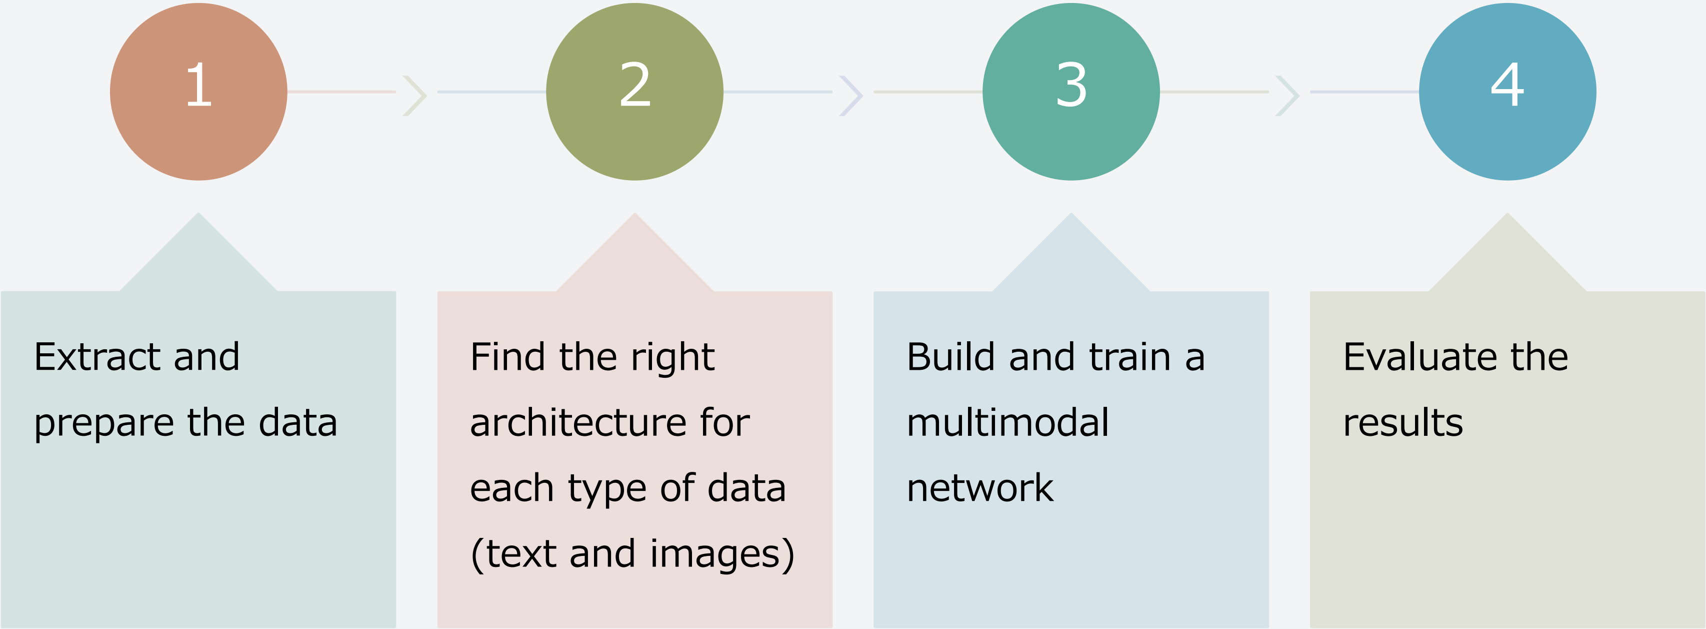

1. Extract and prepare the data
2. Find the right architecture for each type of data (text and images)
3. Build and train a multimodal network
4. Evaluate the results

Dataset

✗ Lack of existing dataset focusing on this kind of task (e.g. Spotify dataset is meant for audio transcription)

Extracted a dataset from Apple Music using their APIs, including nearly 30000 different shows ✓

Each show has metadata regarding title, description, cover image, authors and primary genre

Dataset: cover images

13 Minutes to the Moon (Technology)

The Interior Design Business (Design)

Two Girls Talk Balls (Sports)

Five Point Move - U.S. Greco-Roman Wrestling (Sports)

A Delectable Education Charlotte Mason Podcast (Education)

a delectable
EDUCATION
Spreading the Feast of the Charlotte Mason Method
with Emily Kiser, Liz Cottrill, & Nicole Williams

The Christian Woman Business Podcast (Business)

The Wellness Mama Podcast (Health)

Metabolic Academy (Health)

Quran recitations (Religion)

Dataset: genres

Originally 110 different genres with many overlaps

After pre-processing only 35 different genres with a rather unbalanced distribution

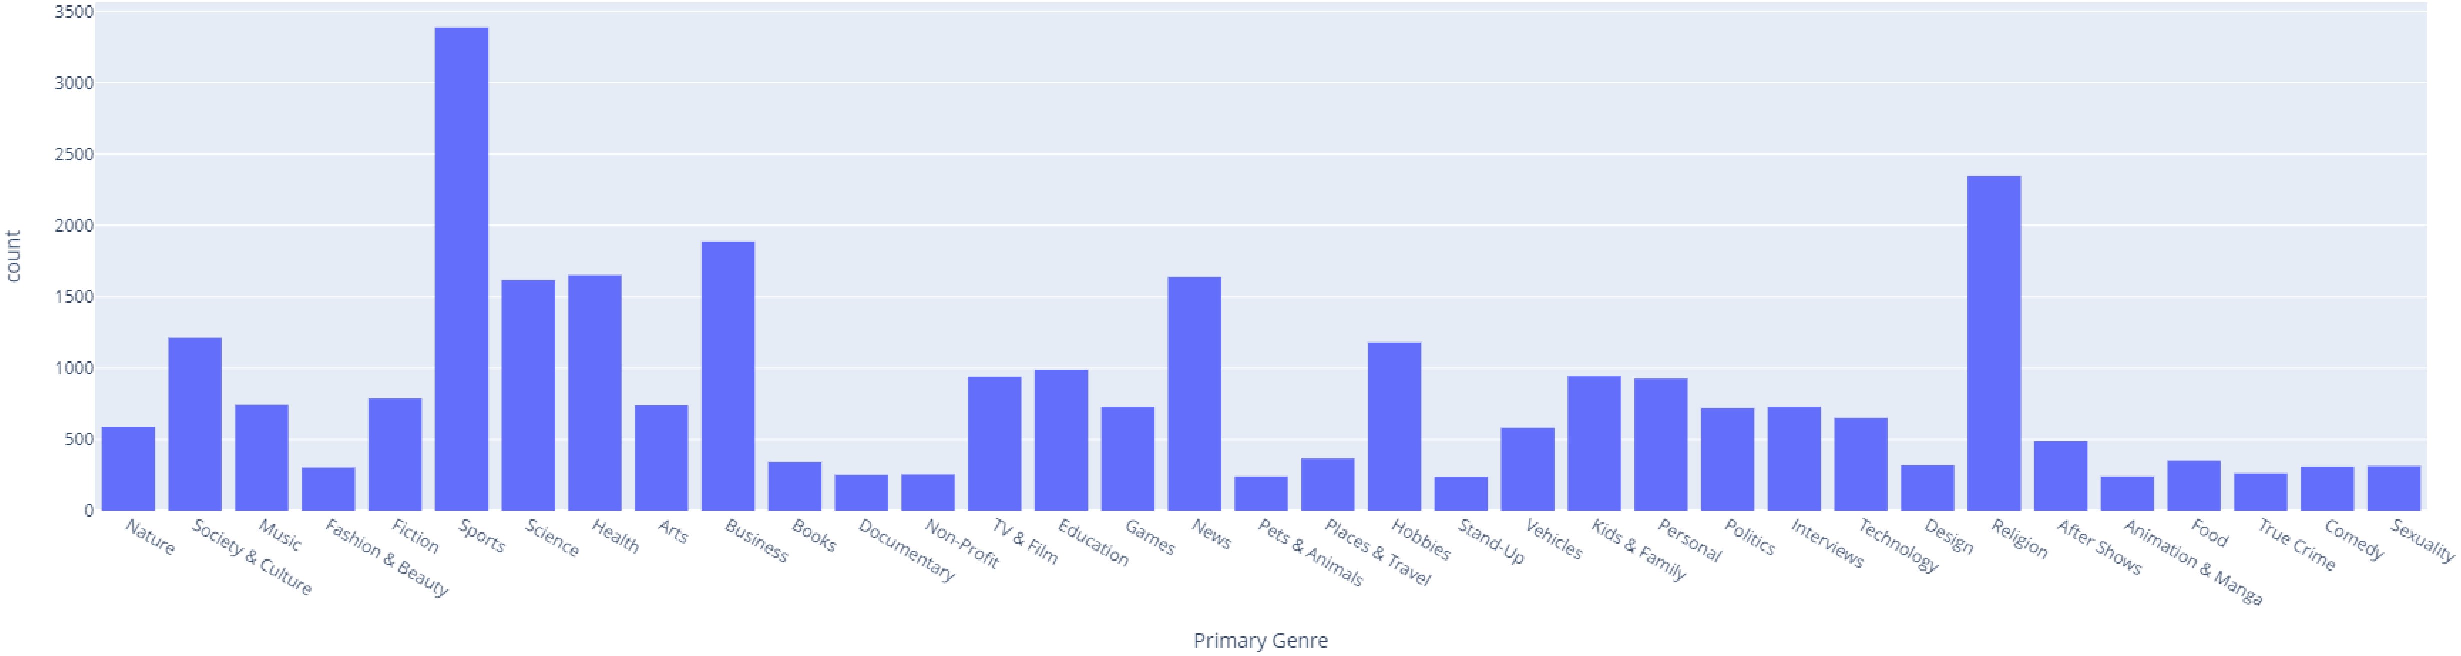

Dataset: descriptions

Two Girls Talk Balls (Sports)

Title: Two Girls Talk Balls

Description: «An alternative take on Women's Football. Covering the WSL, NWSL, Euro 2022 and Women's Champions League; Charlotte and Tamsin bring a refreshing voice with debate, banter and highlights from the best of the Women's Game. Proudly sponsored by FotMob”

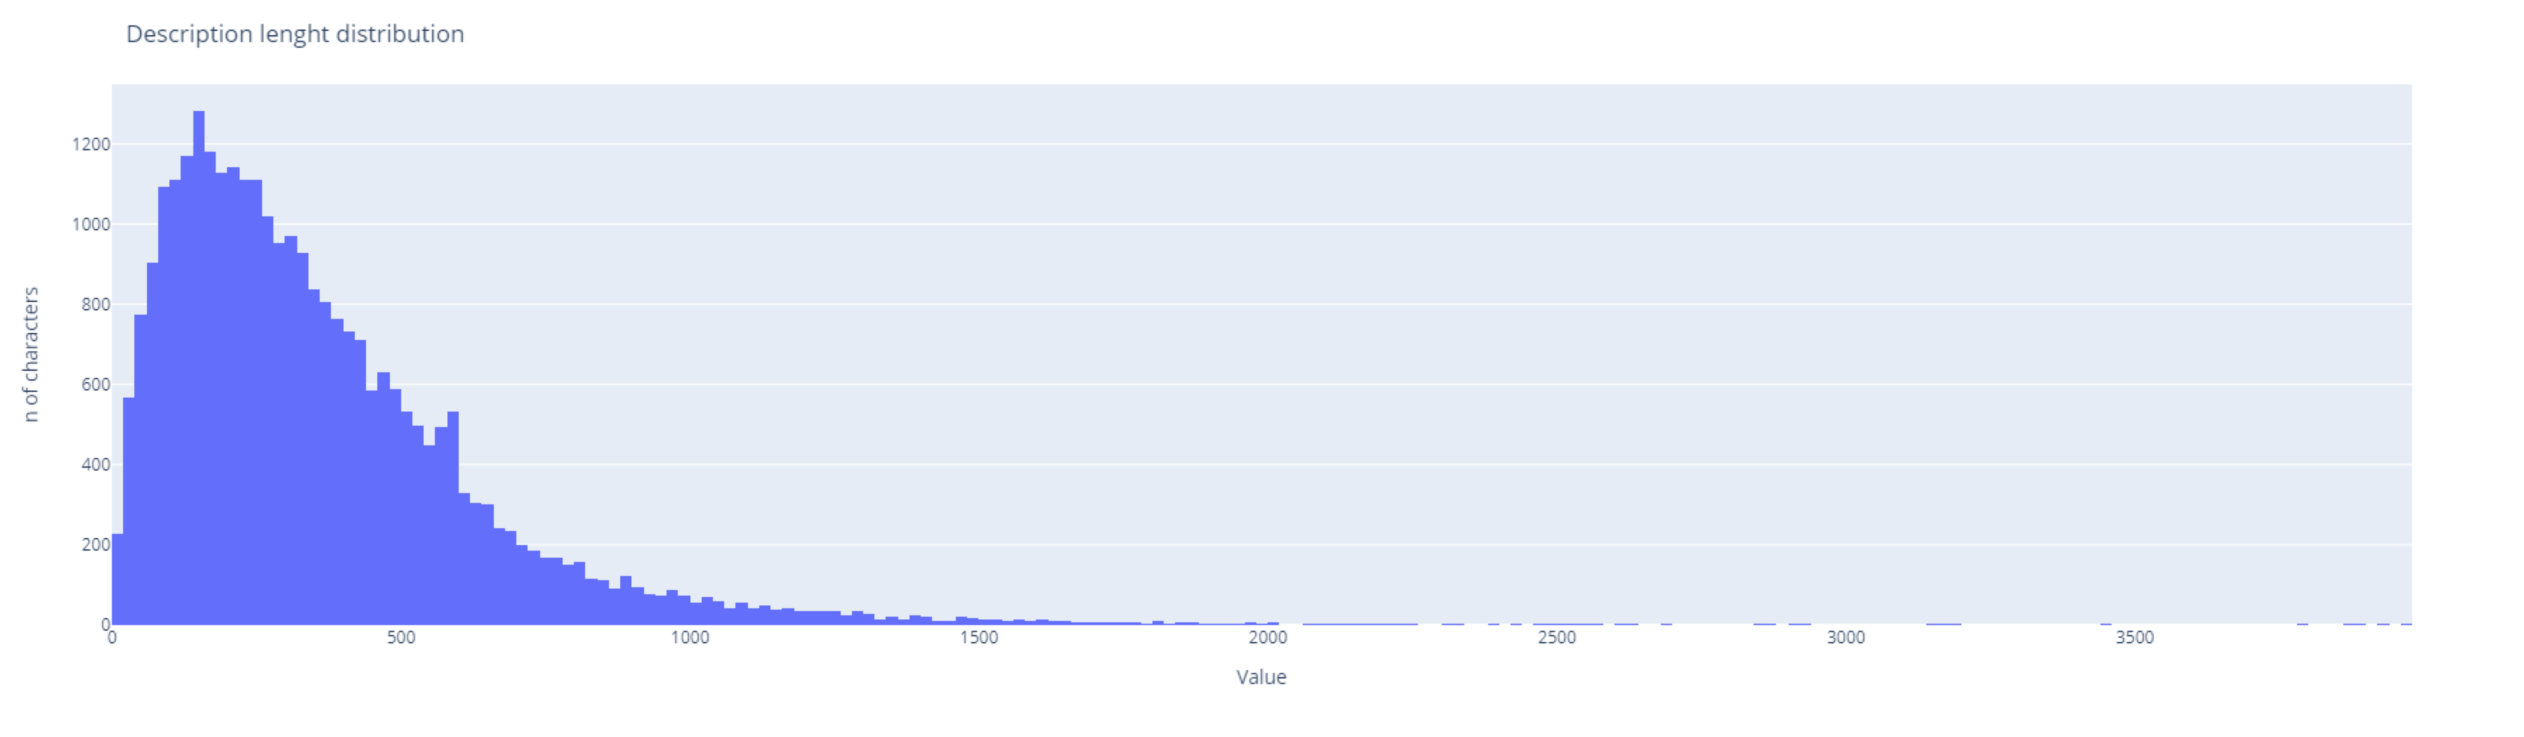

DNN: images

Images are coded as a numerical matrix (height x width x channels)

Each image in the dataset is in JPEG format and has a shape of 600 x 600 x 3 (RGB)

DNN can learn how to filter (parts of) an image to extract relevant features (Convolutional Neural Network)

Instead of building from scratch a CNN, better performances are achieved use a pre-trained architecture (transfer learning)

VGG16

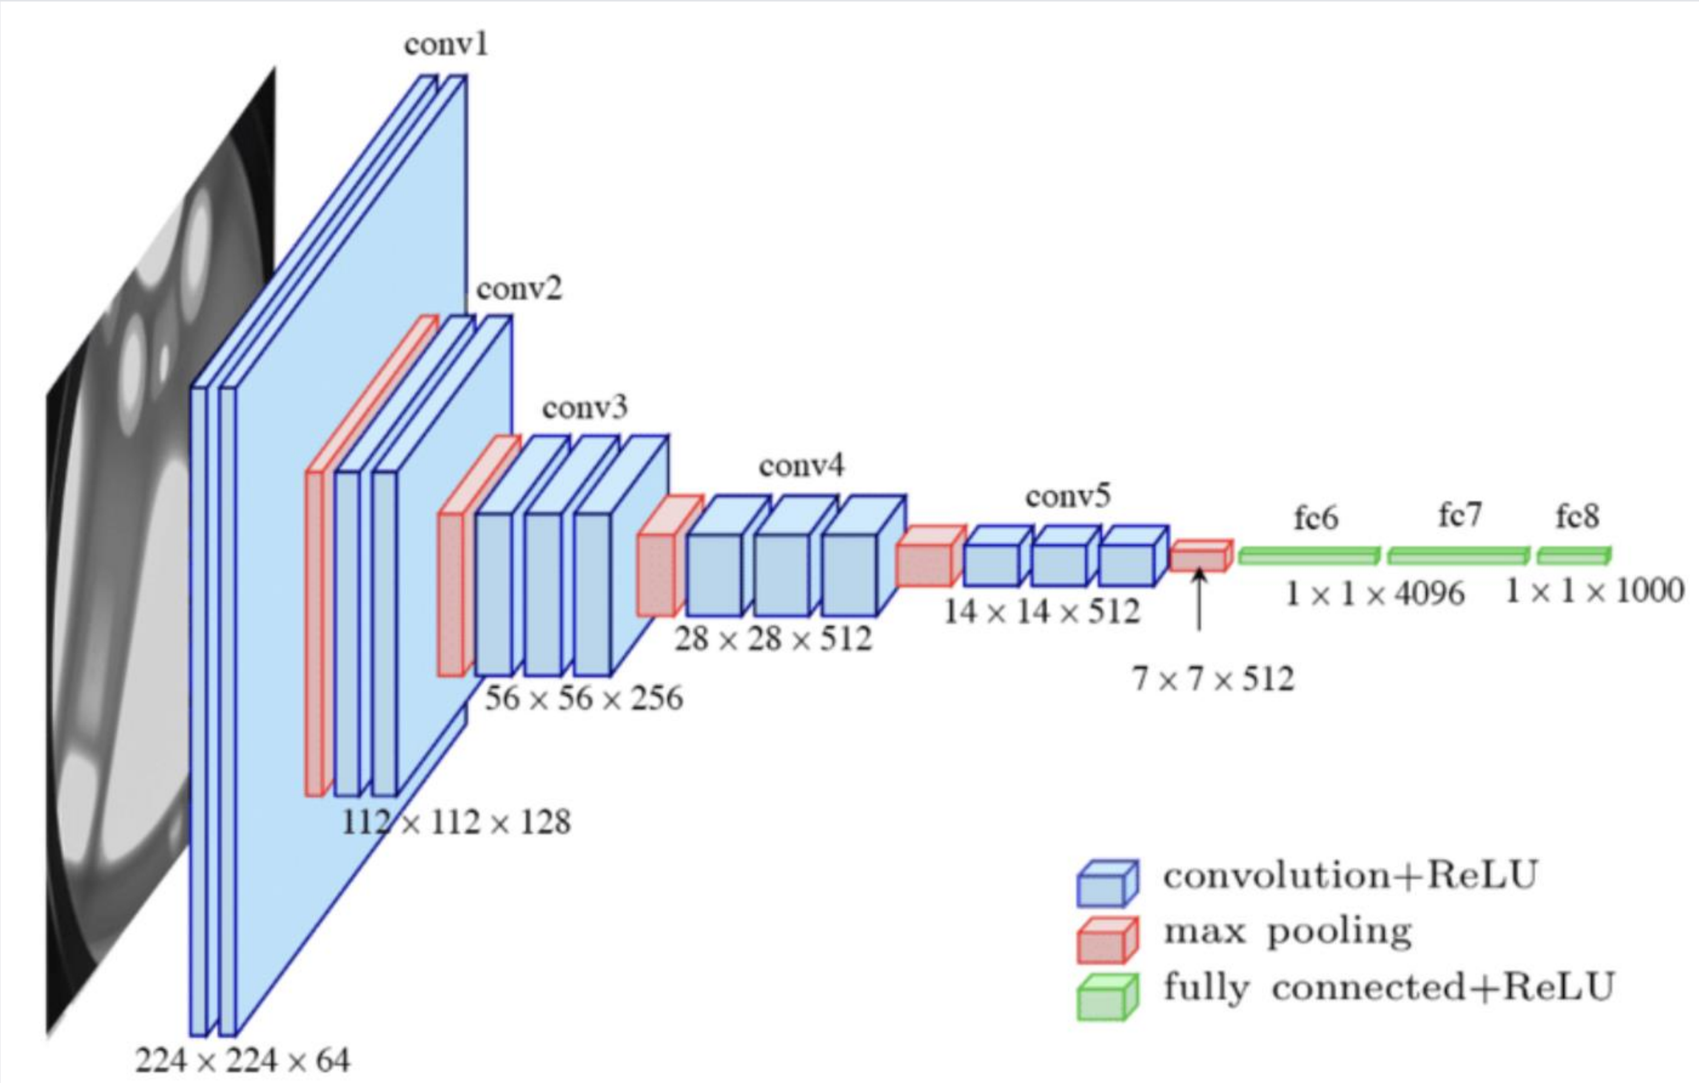

Large architecture trained on millions of images (ImageNet dataset)

Stack of convolutional layers and max pooling using ReLU as activation function

It can be used for different tasks without re-training thanks to transfer learning

VGG16

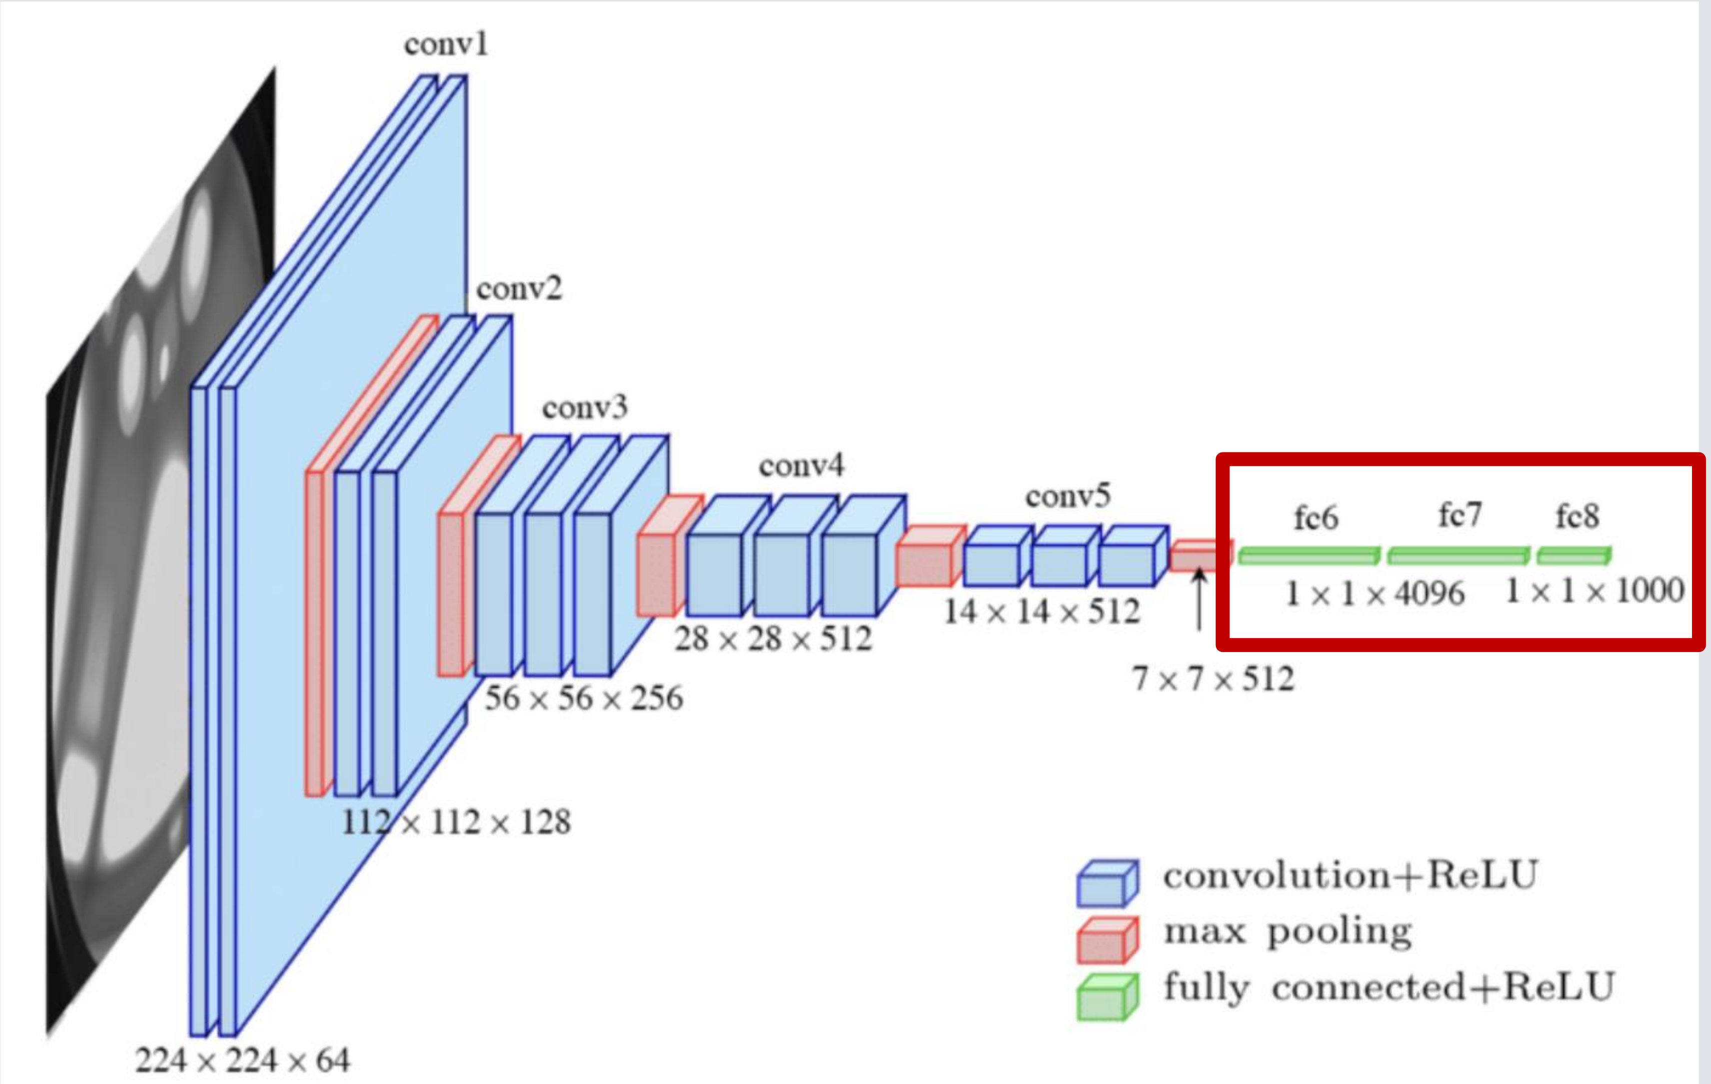

Cut the last layer (classifier head) and replace it with a dense layer with a neuron for each label (35)

Reshape images to 224 x 224 size and convert them from RGB to BGR (automatically pre-processed by Keras)

Use data augmentation to create artificially additional samples for training (random flip and random rotation)

DNN: descriptions

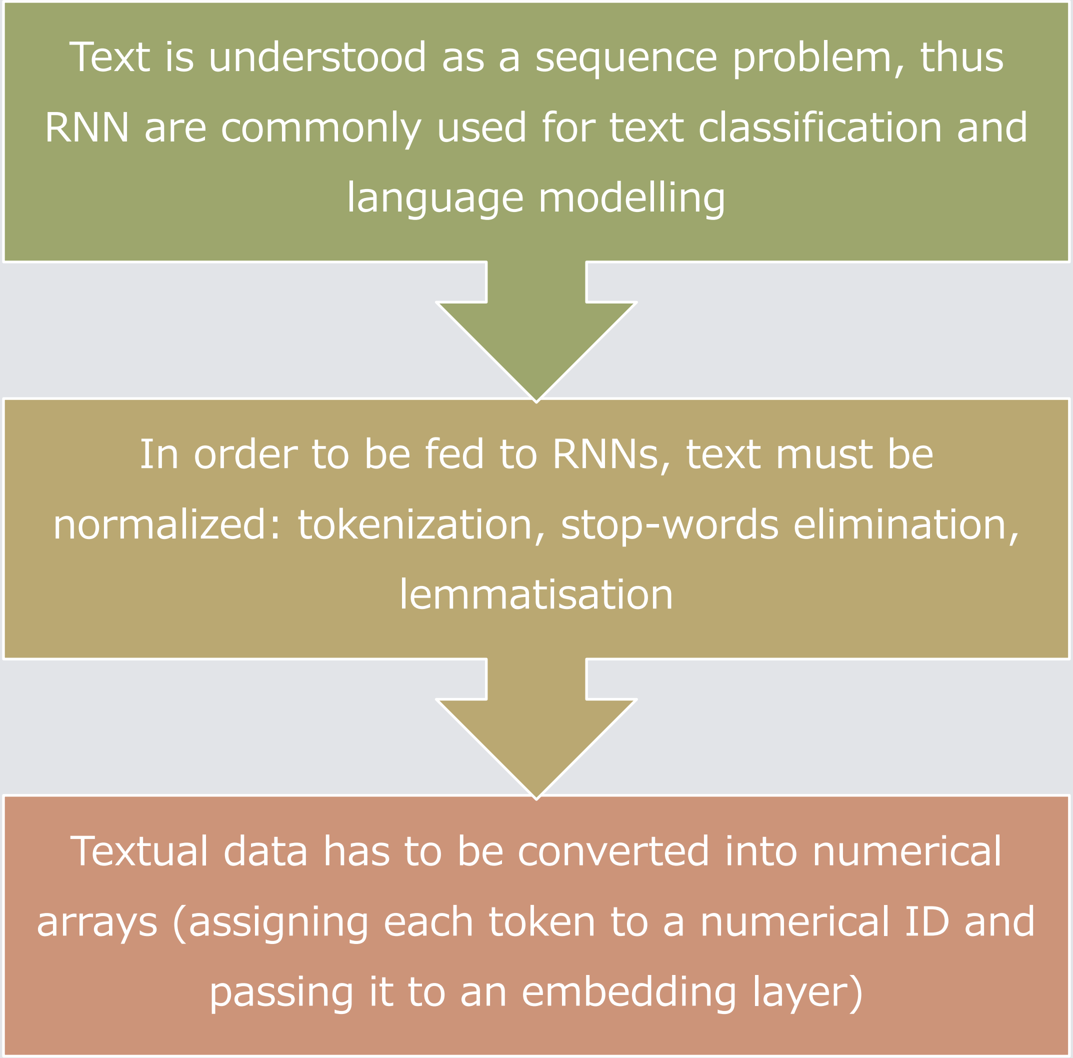

Text is understood as a sequence problem, thus RNN are commonly used for text classification and language modelling

In order to be fed to RNNs, text must be normalized: tokenization, stop-words elimination, lemmatisation

Textual data has to be converted into numerical arrays (assigning each token to a numerical ID and passing it to an embedding layer)

DNN for text: GRU

Standard RNNs have problems with long sequences, which is rather usual in textual data (vanishing gradient problem)

LSTM are able to solve this problem by implementing a form of long-term memory

GRU is a special kind of LSTM, faster to train and (usually) more efficient

Brief introduction to LSTM

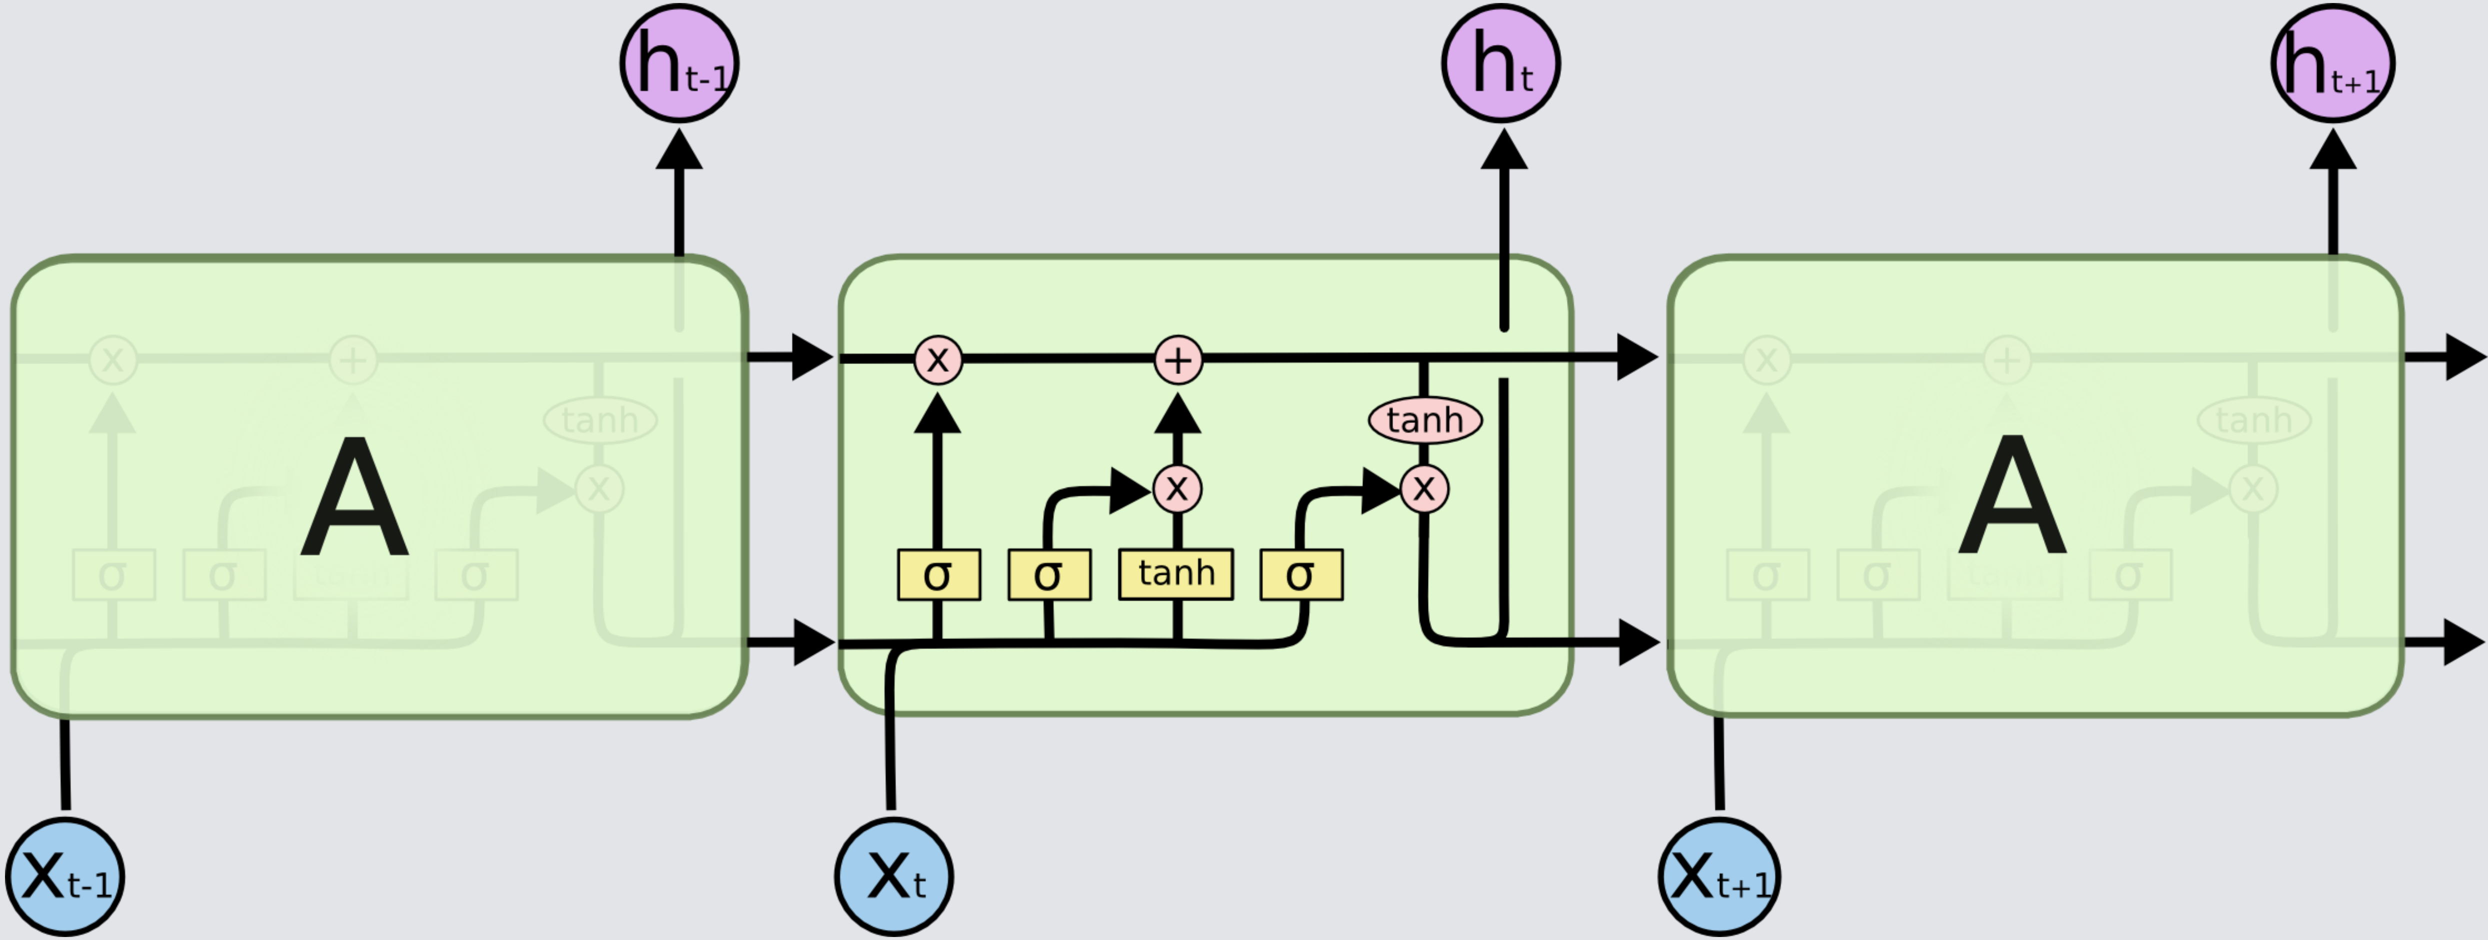

GRU

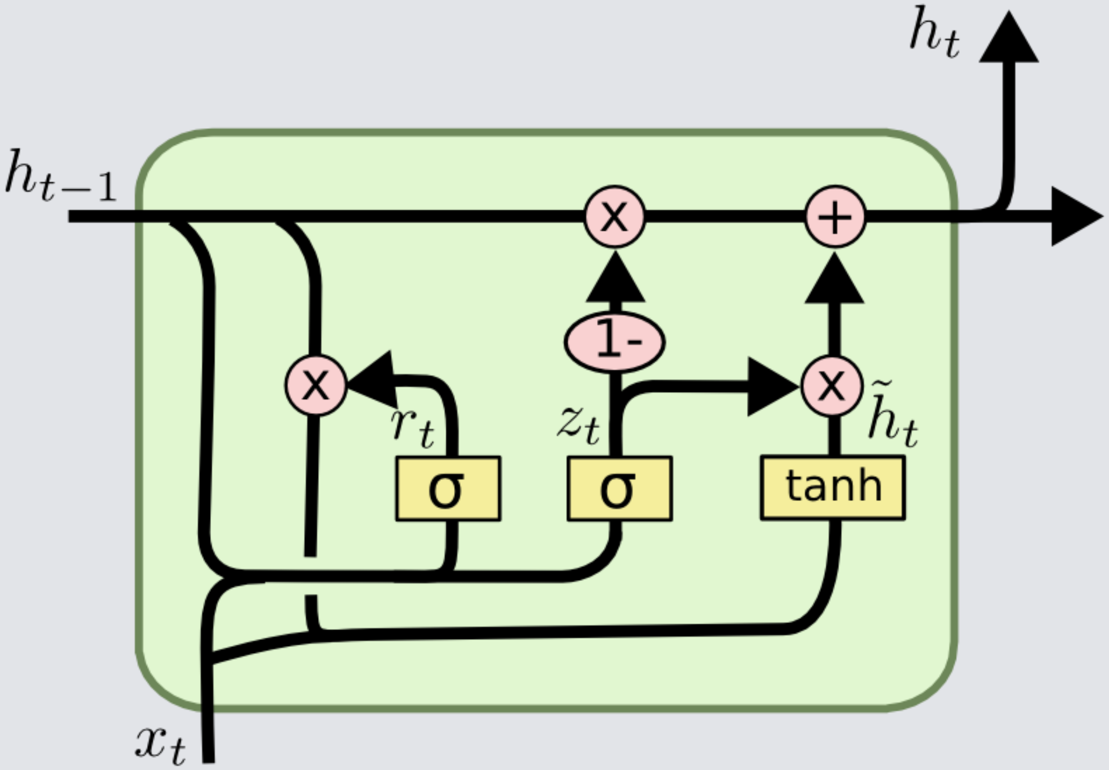

Simpler than normal LSTM

Combines forget and input gate into a single «update layer»

Merges cell states and hidden states into one

Reset gate decides whether the previous cell is important or not

Multimodal network

Data fusion

Framework in which multiple data source are mixed to make them more useful for a specific application

It can mix information to extract more knowledge and have a better form of abstraction of the problem

The goal is to improve the system performances both in terms of accuracy and robustness

Particularly useful in robotics or computer vision in general

Late vs early fusion

In the literature the distinction between early and late fusion depends on the position of the fusion process

Late: it fuses the results of two or more different computational streams that work in parallel (this case)

Early: it merges data sources at the beginning of the process to create a joint representation of the data

Keras implementation

The project was implemented with Keras high level APIs, splitting the original dataset into training, validation and test sets

For VGG16 and the GRU model the standard Generator class was used to feed the network batches of data without saturating memory

The multimodal network required the development of a custom generator in order to feed mixed type of data to the network

Training was performed on Google Colab and Amazon Sagemaker Studio Lab (Tesla K80 vs Tesla T4)

Standard hyperparameters (Adam optimizer, batch size of 32, categorical cross-entropy as loss function) and early-stopping to avoid excessive overfitting

The multimodal network (VGG16 + GRU)

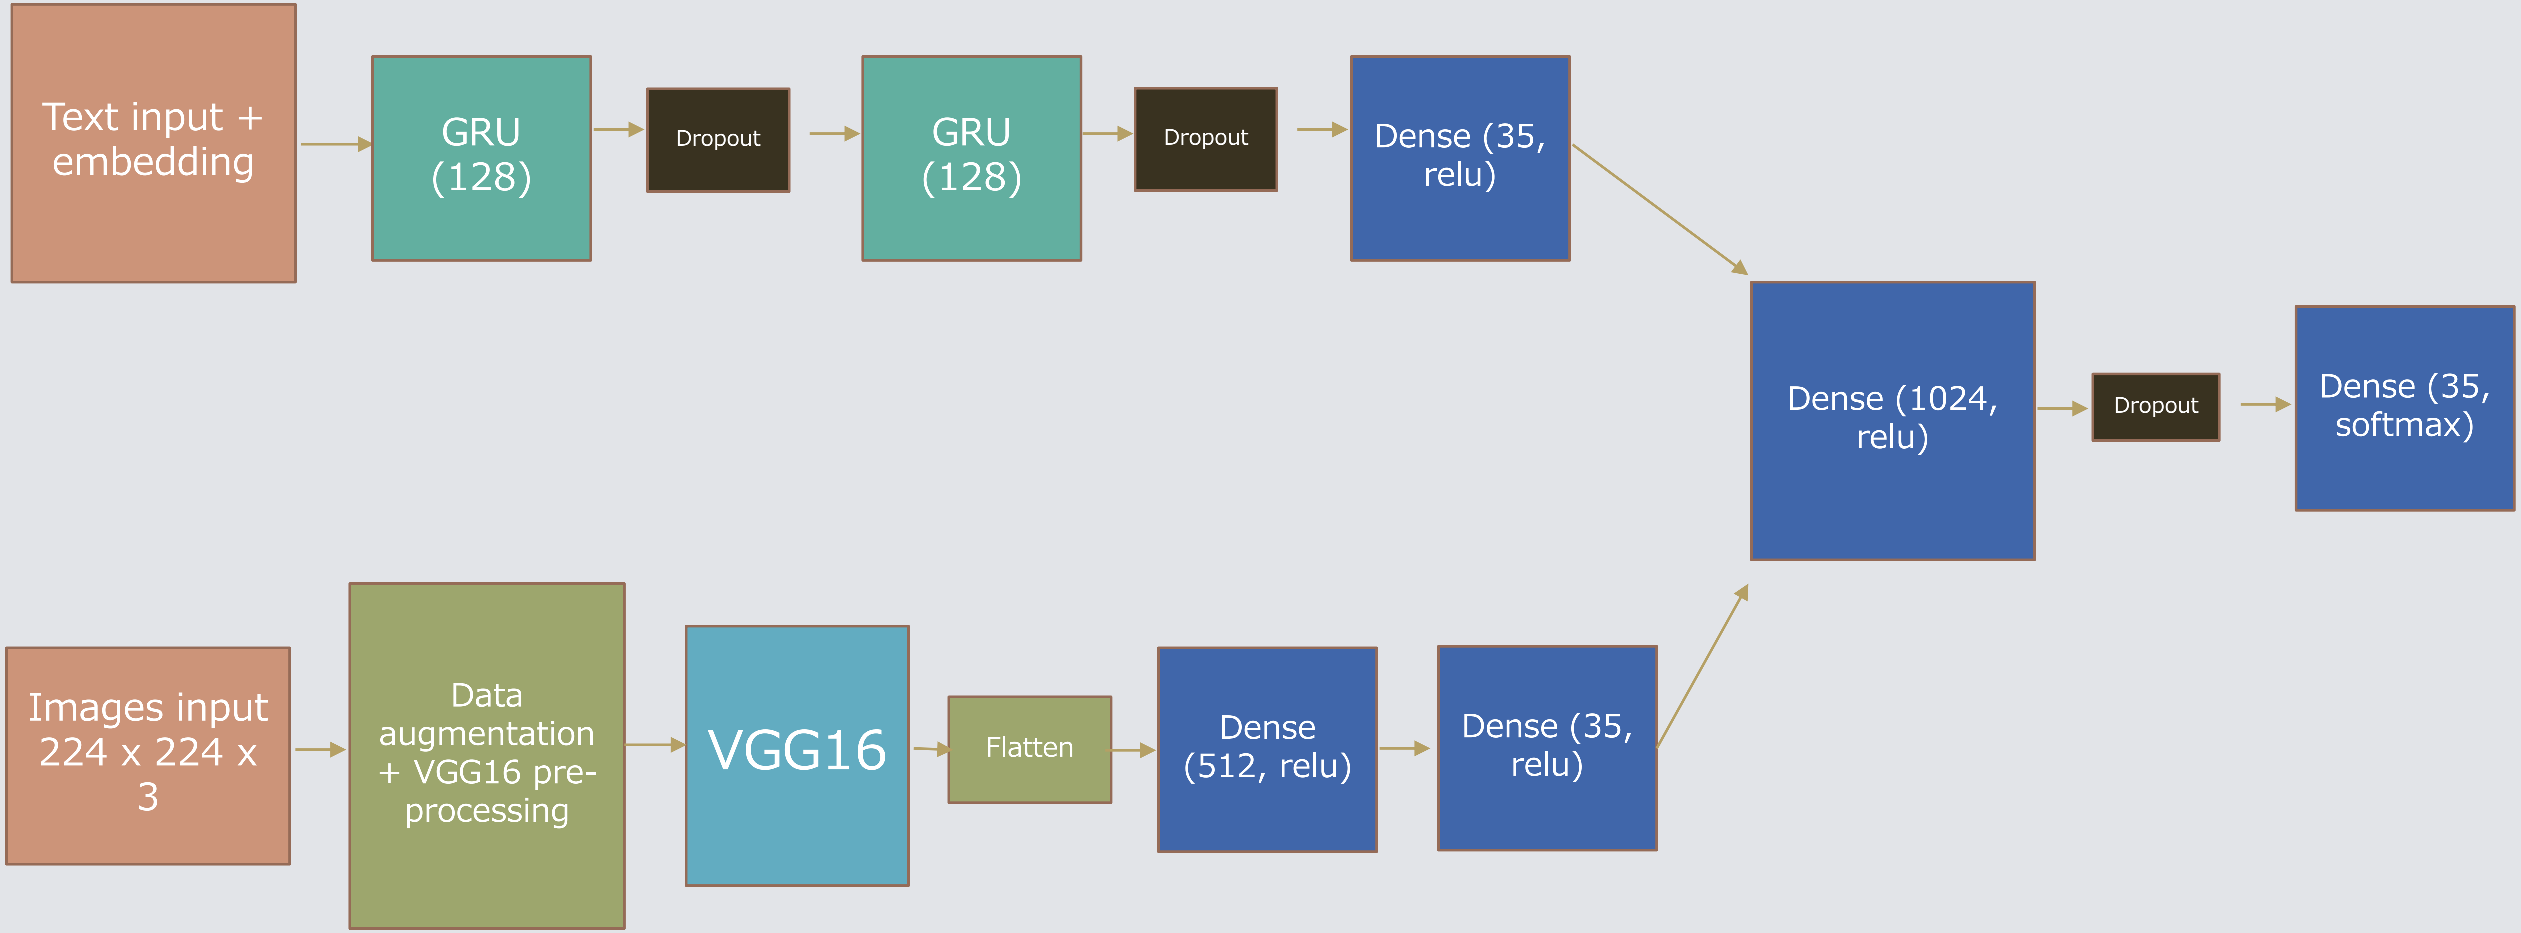

Results

Compared losses

Compared loss

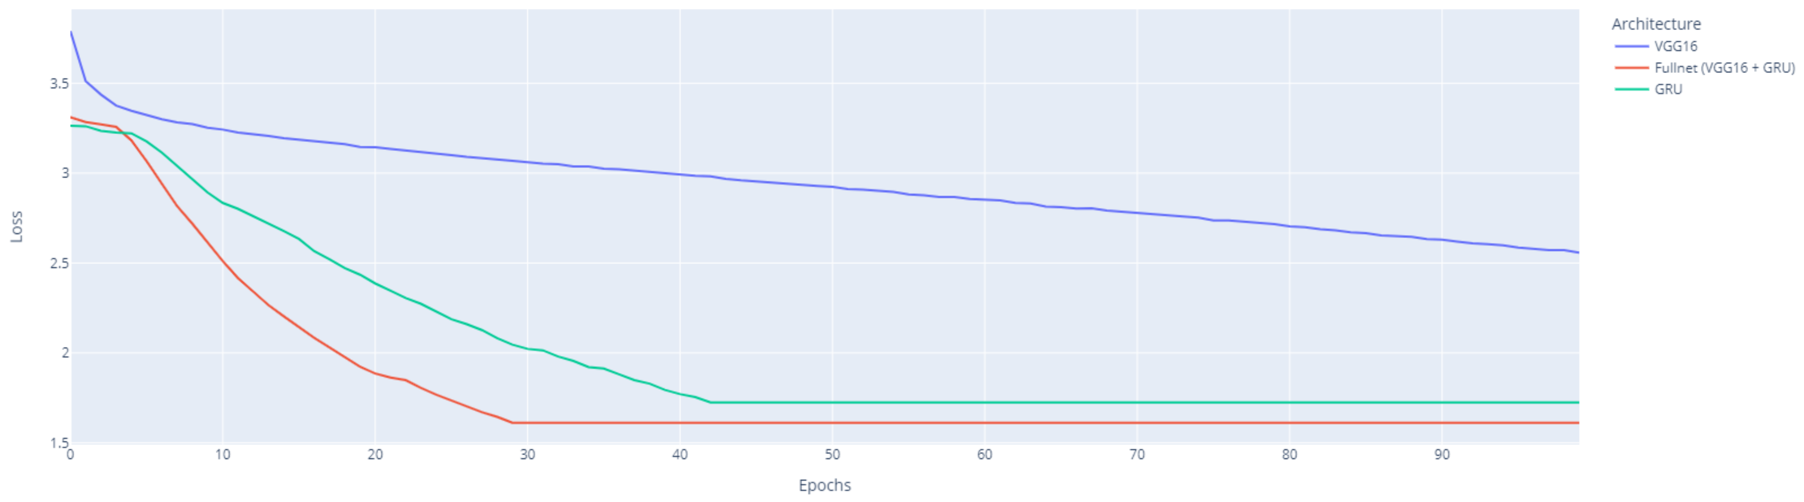

Compared validation loss

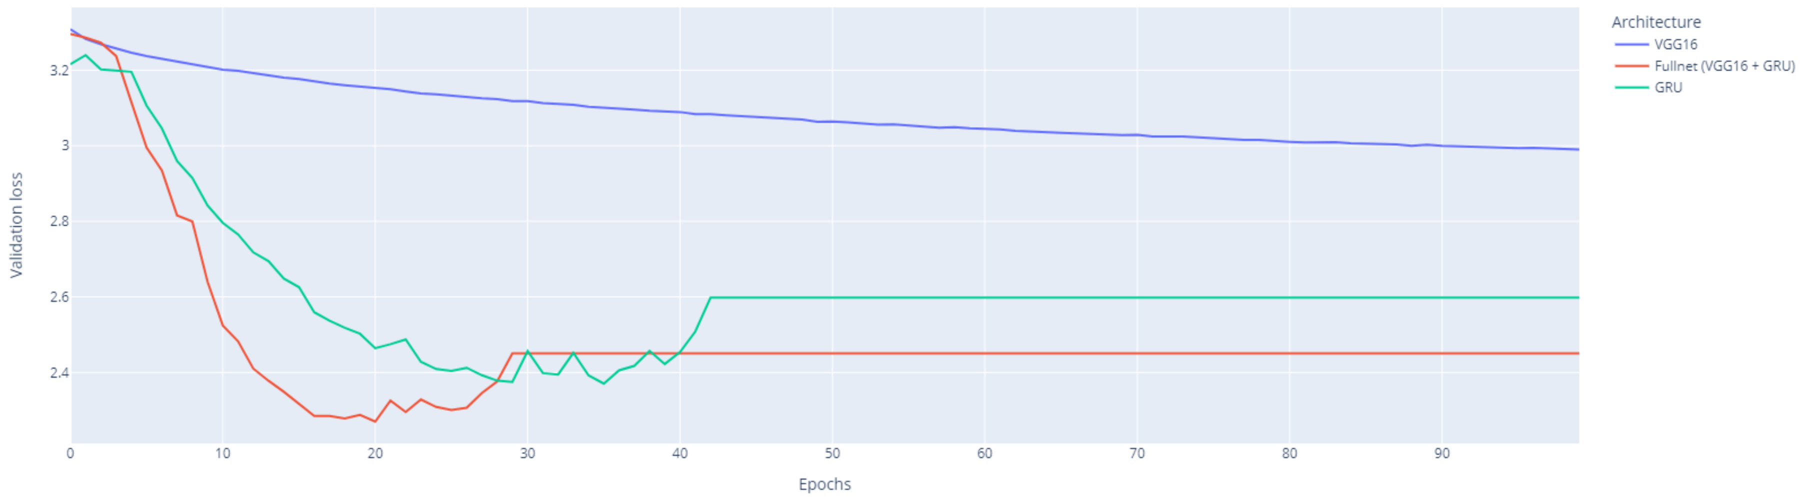

VGG16 required a much slower learning rate to gain adequate performances

GRU and multimodal network are much faster and able to reach lower loss both in training and validation

Early stopping prevents excessive overfitting

Compared accuracy

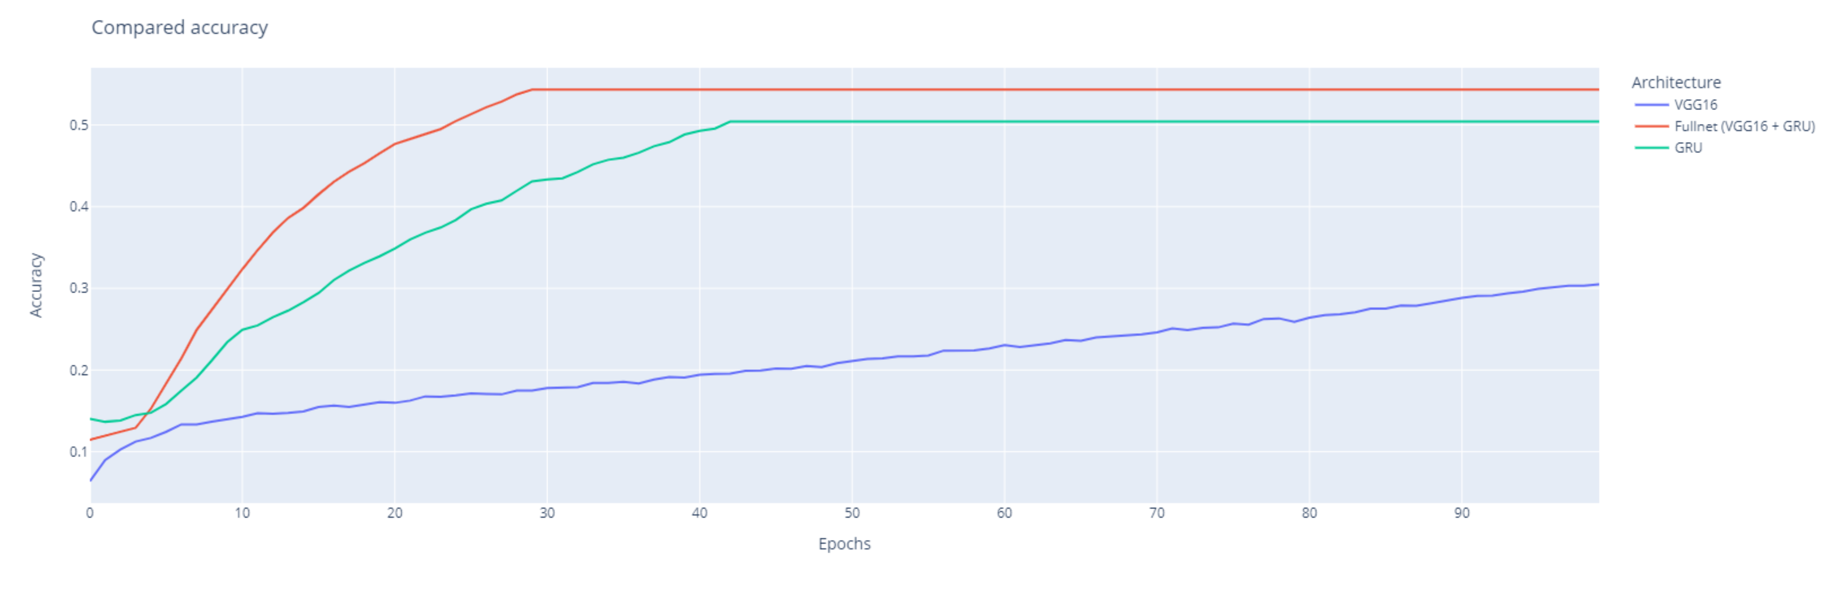
Compared accuracy

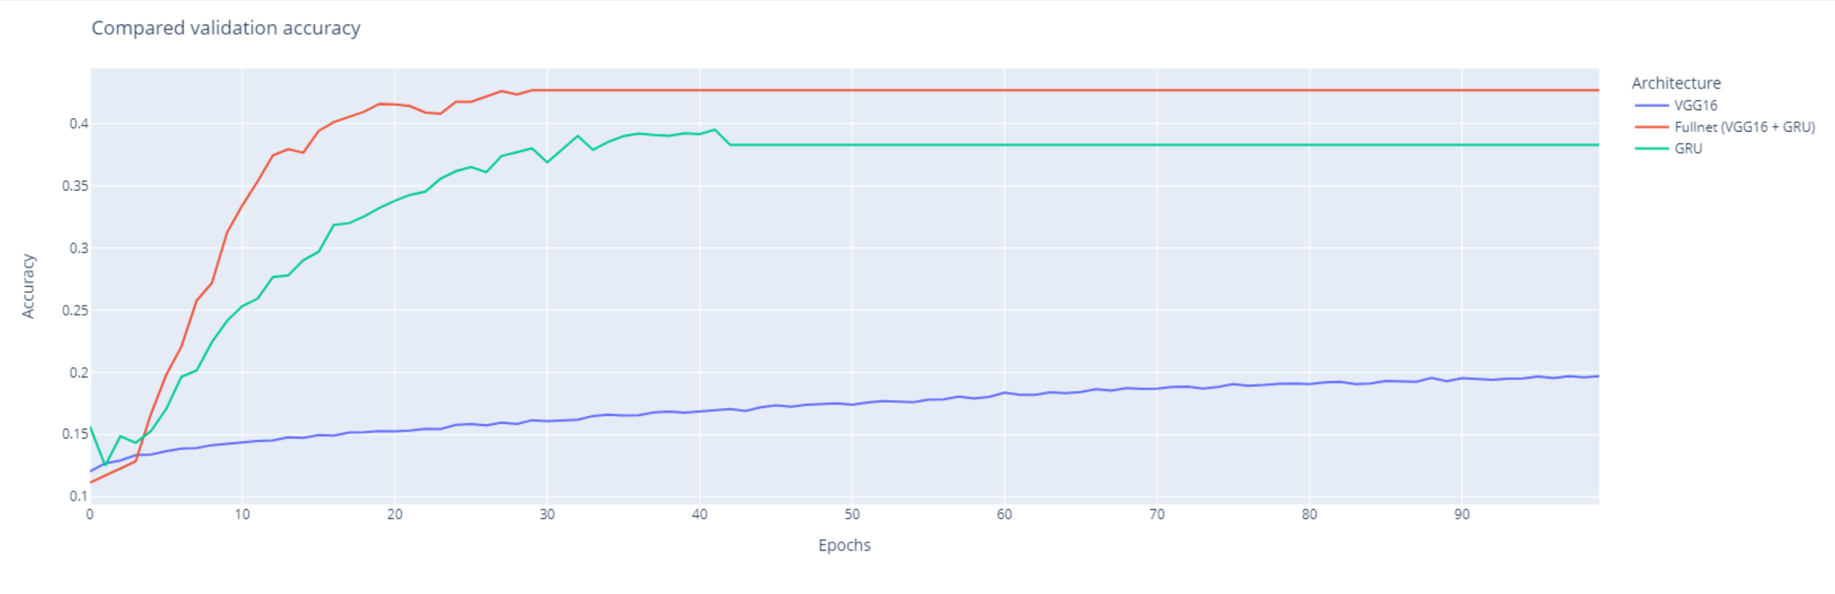
Compared validation accuracy

VGG16 learns very slowly over time

GRU and multimodal network are much faster and able to reach lower loss both in training and validation

Early stopping prevents excessive overfitting

Evaluation metrics on testing set

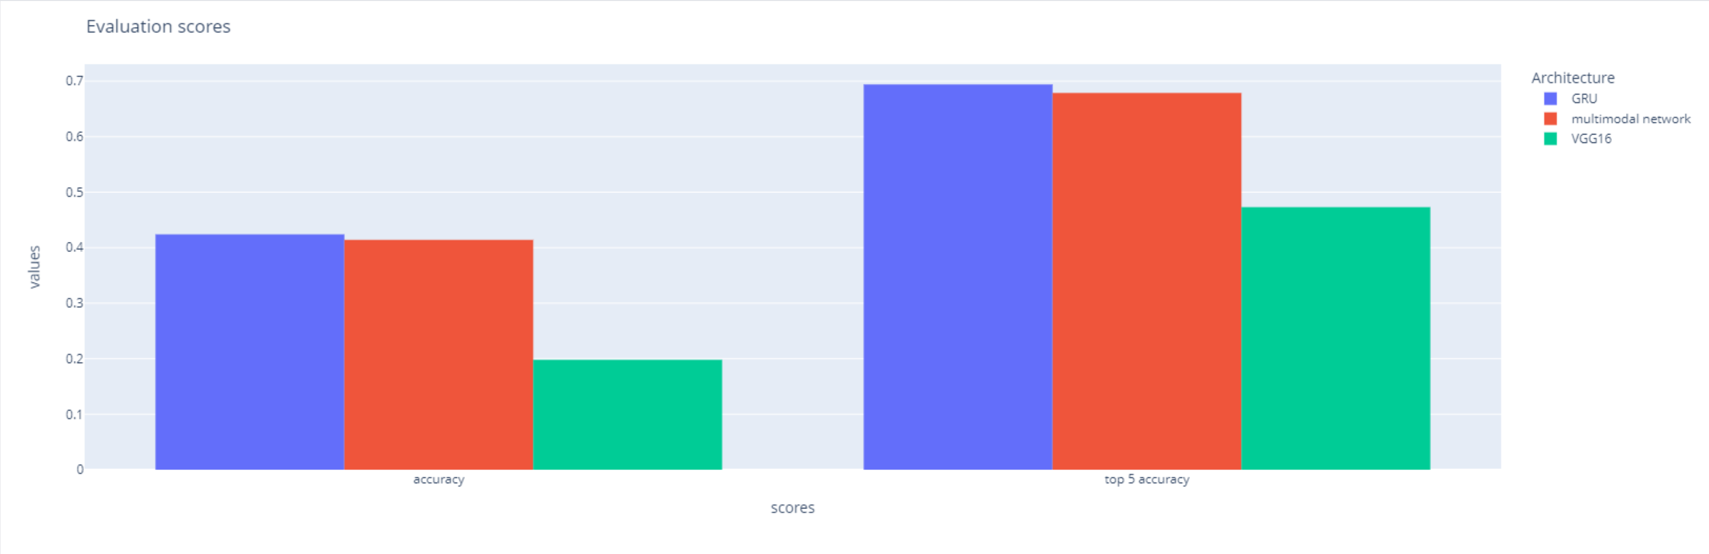

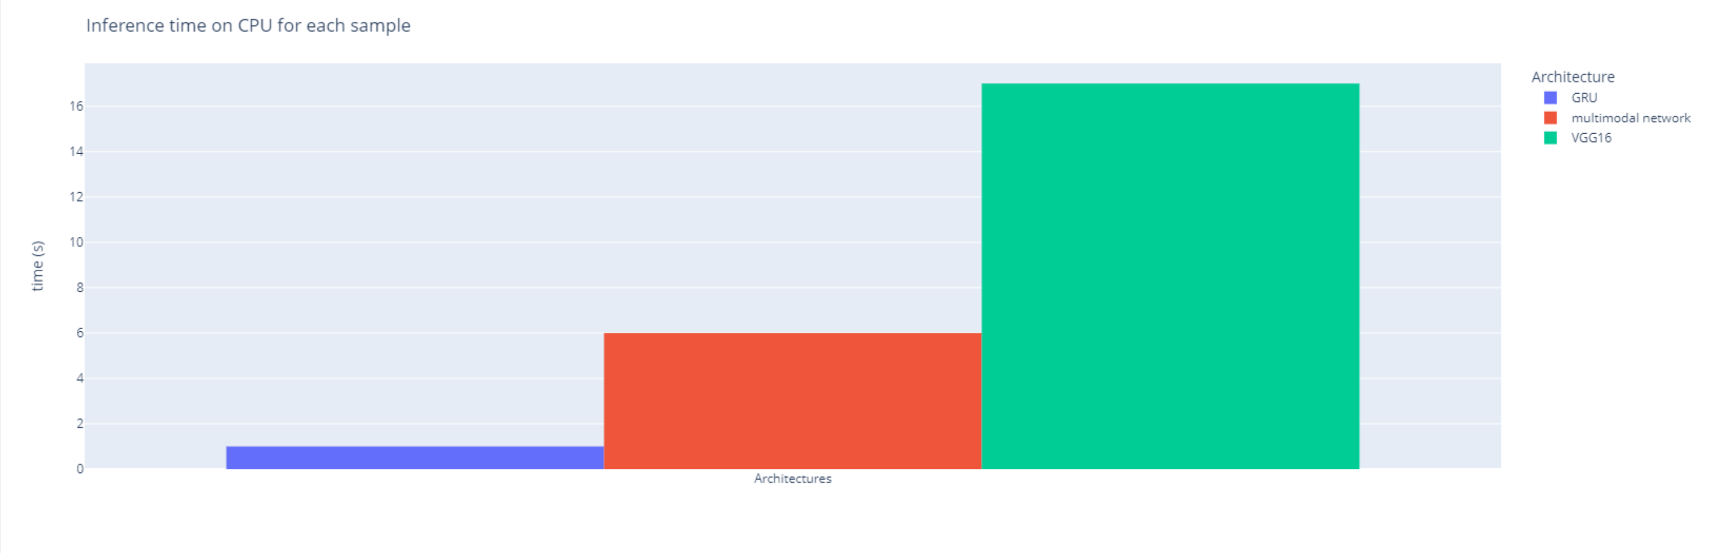

GRU is slightly more accurate on the testing set

If we consider top 5 accuracy the difference between GRU/multimodal and VGG16 is less significant

Inference time increase for the multimodal network, but not as dramatically expected looking at how VGG16 performs on its own

Predictions: VGG16

Predicted: Sports
Real: News

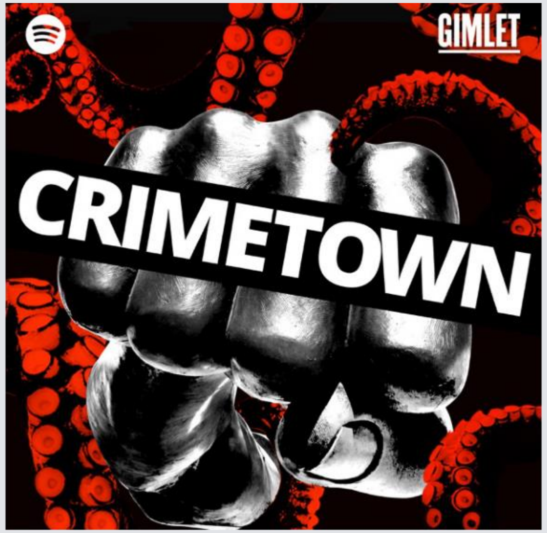

Predicted: Religion
Real: Vehicles

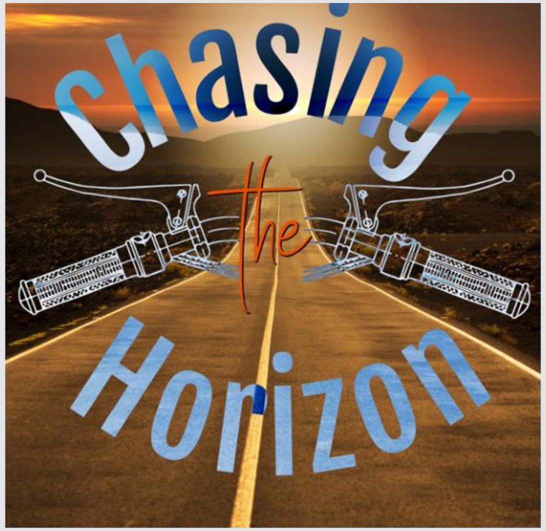

Predicted: Society & Culture
Real: Religion

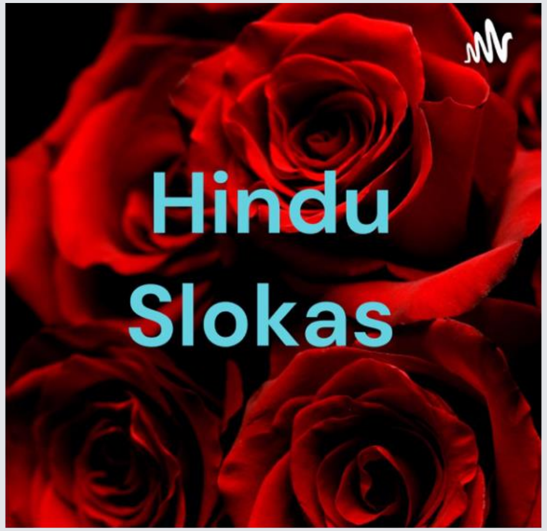

Predicted: Sports
Real: Documentary

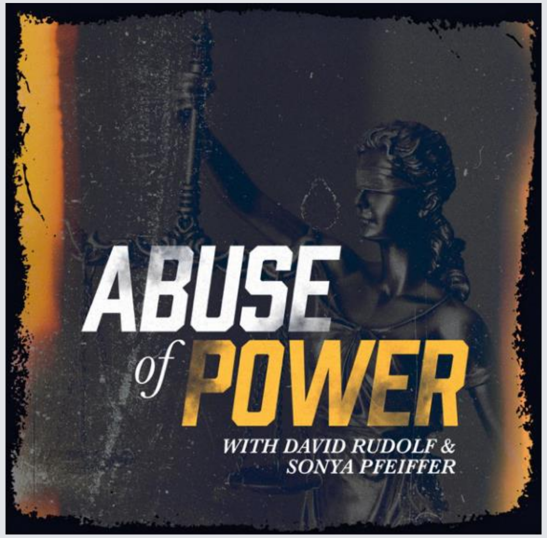

Predicted: Business
Real: Religion

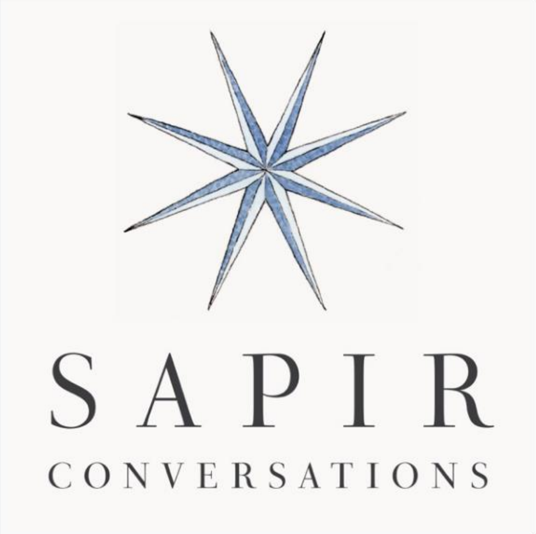

Predicted: Sports
Real: Society & Culture

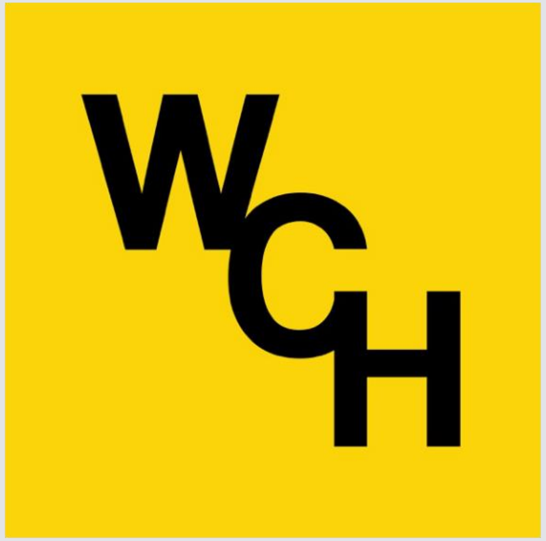

Predictions: GRU

Predicted: Sports
Real: News

GIMLET
CRIMETOWN

Predicted: Religion
Real: Vehicles

Chasing the Horizon

Predicted: Religion
Real: Religion

Hindu Slokas

Predicted: Interviews
Real: Documentary

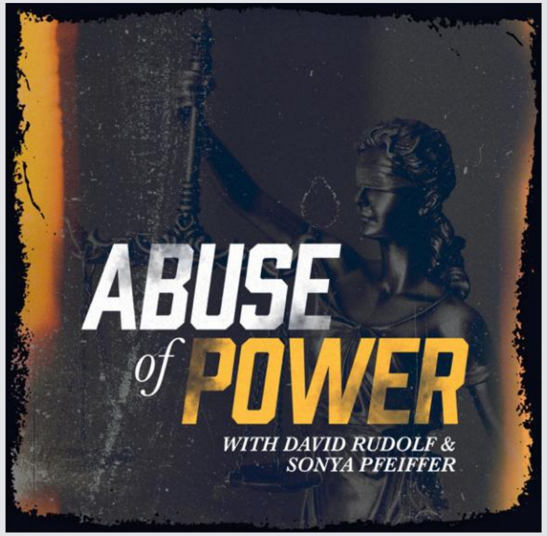

Predicted: Fiction
Real: Religion

SAPIR
CONVERSATIONS

Predicted: TV & Film
Real: Society & Culture

Predictions: multimodal network

Predicted: True Crime
Real: News

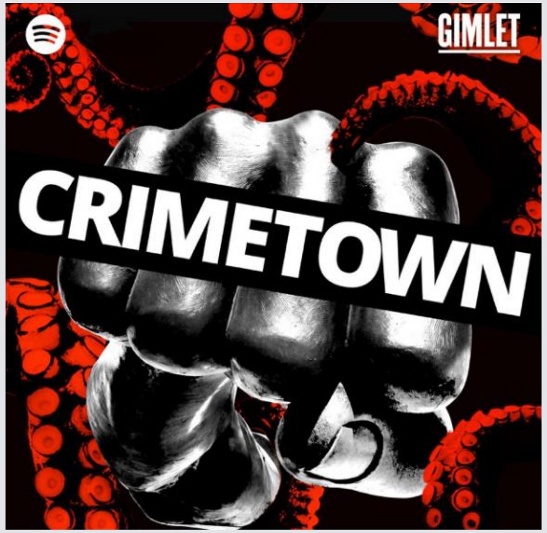

Predicted: Vehicles
Real: Vehicles

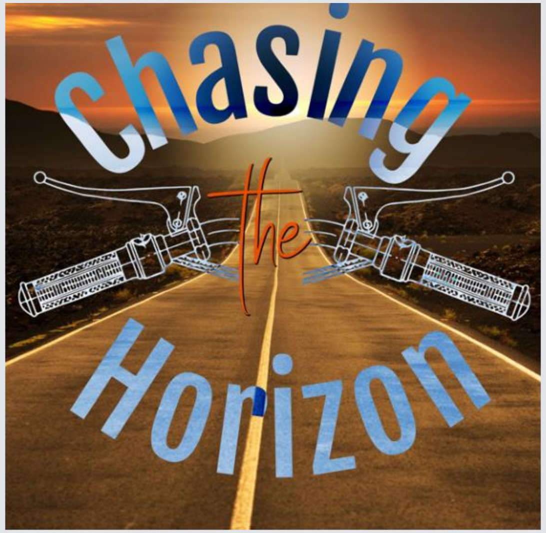

Predicted: Vehicles
Real: Religion

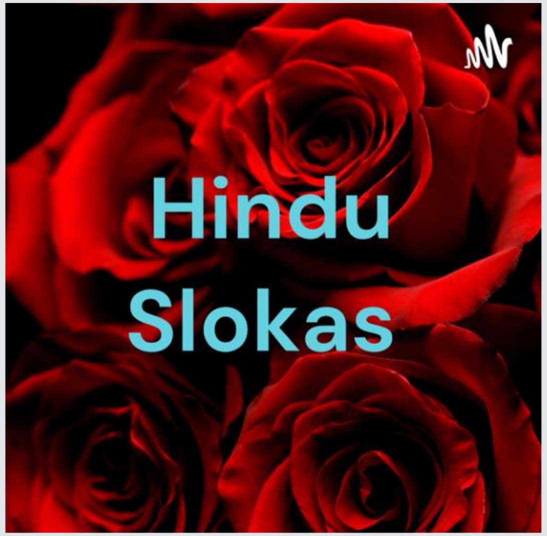

Predicted: Arts
Real: Documentary

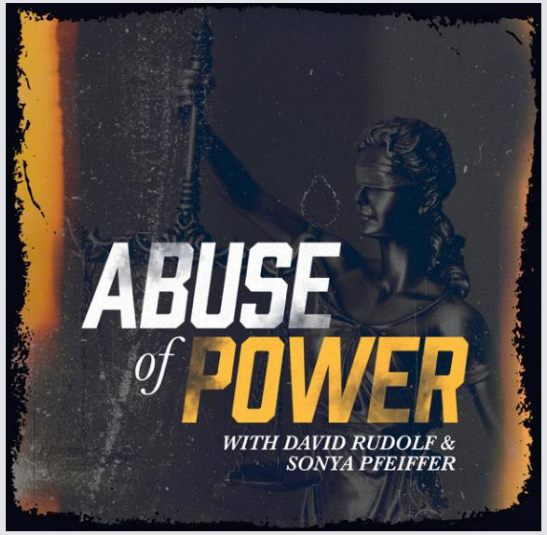

Predicted: Religion
Real: Religion

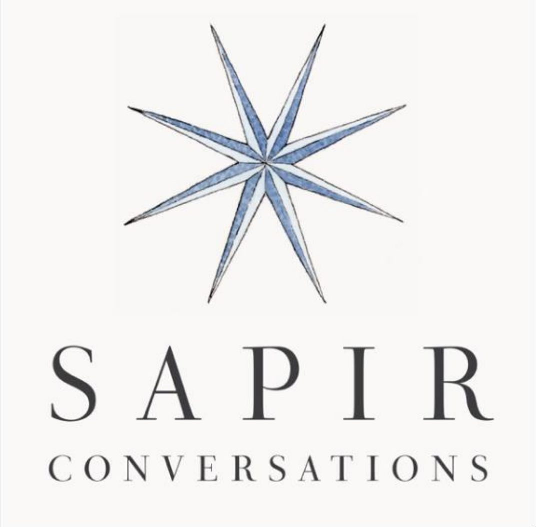

Predicted: Politics
Real: Society & Culture

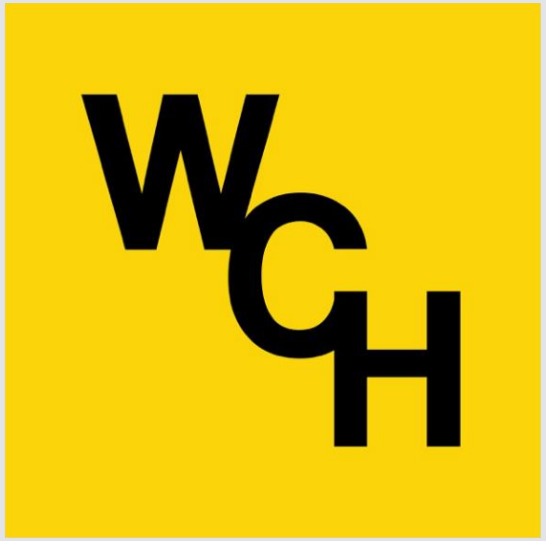

Balance vs unbalanced metadata

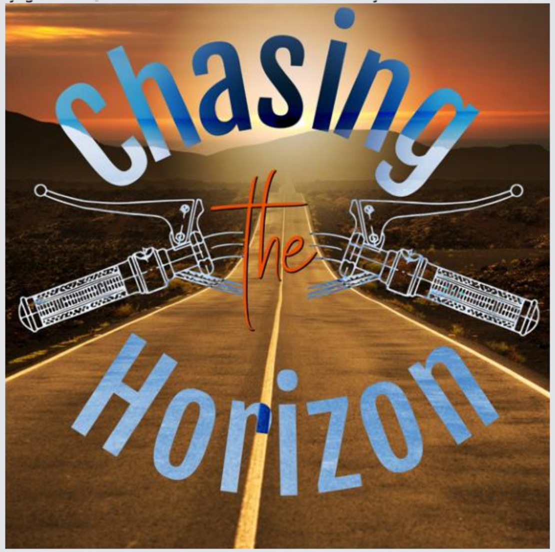

«Chasing the Horizon is a podcast by, about and for motorcyclists. We talk to motorcycle industry figures, technical experts and riders just like you no matter what bike they love to ride. Please sign up on the mailing list at http://tinyletter.com/chasingthehorizon and subscribe in popular podcast apps or listen to every episode at http://chasingthehorizon.us»

VGG16: Religion
GRU: Religion
Multimodal: Vehicles

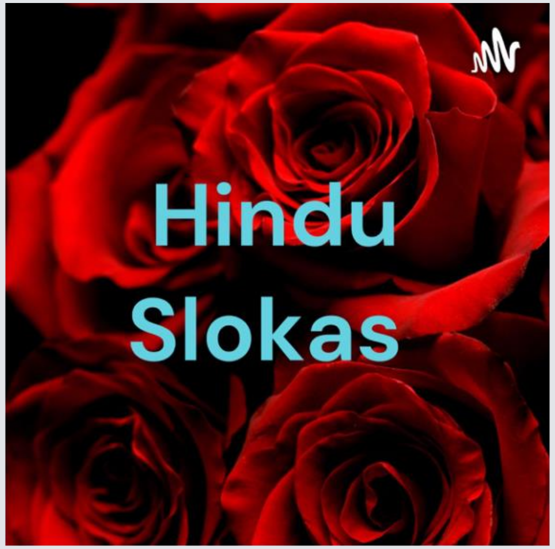

«Hindu slokas»

VGG16: Society & Culture
GRU: Religion
Multimodal: Vehicles

Conclusions

The multimodal model converges faster (in terms of epochs) during training and it tends to have better accuracy

However, the GRU model performs a little better on the testing set: this might be due to the fact that the dataset is largely unbalanced and so is the testing set

The multimodal approach seems to help in cases in which the two features are likely mixed, but it might lead to mistakes in unbalanced cases (e.g. less text and confusing image)

Labels were manually and arbitrarily aggregated; we might get better results depending on how genres were merged

Further experiments with larger and more balanced dataset could improved performances

Bibliography

Miller, Stuart J., et al, «Multi-Modal Classification Using Images and Text», SMU Data Science Review 3.3 (2020): 6.

Gadzicki, Konrad, et al., «Early vs late fusion in multimodal convolutional neural networks», 2020 IEEE 23rd International Conference on Information Fusion (FUSION), IEEE, 2020.

Li, Jonathan, Di Sun, and Tongxin Cai, «Genre Classification via Album Cover».

Ofli, Ferda, Firoj Alam, and Muhammad Imran, «Analysis of social media data using multimodal deep learning for disaster response», arXiv preprint arXiv:2004.11838 (2020).

Bibliography

Yin, Wenpeng, et al., «Comparative study of CNN and RNN for natural language processing», arXiv preprint arXiv:1702.01923 (2017).

Ahmed, Eman, and Mohamed Moustafa, «House price estimation from visual and textual features», arXiv preprint arXiv:1609.08399 (2016).

Nawaz, Shah, et al., «Are these birds similar: Learning branched networks for fine-grained representations», 2019 International Conference on Image and Vision Computing New Zealand (IVCNZ), IEEE, 2019.